

The Deep Latent Position Block Model for Clustering and Representation of Networks

Rémi Boutin¹, Pierre Latouche² and Charles Bouveyron³

¹ LPSM - Sorbonne Université

² LMBP - Université Clermont Auvergne

³ Maasai team - INRIA, Université Côte d'Azur


Journées de la Statistique, Bordeaux, May 2024




Introduction and motivation

The **networks** are a natural data structure to represent interactions between objects or individuals, such as:

- ▶ emails, co-authorship networks
- ▶ biological networks (protein-protein interactions networks)
- ▶ social websites (Facebook, Twitter)



(a) Cora network.



(b) Example of an enzyme network.

Figure: These networks representations were computed with the Fruchterman-Reingold algorithm¹.

¹fruchterman1991graph.

Outline

1. Introduction
2. Generative model
3. Inference
4. Evaluation on synthetic data
5. Analysis of a real world dataset
6. Conclusion

Stochastic Block Model

The stochastic block model² assumes that each node is assigned to a single cluster:

$$\boldsymbol{\eta}_i \stackrel{i.i.d}{\sim} \text{Multinomial}(1; \alpha = (\alpha_1, \dots, \alpha_d)). \quad (1)$$

Hence,

$$\eta_{ik} = \begin{cases} 1 & \text{if } i \text{ is in cluster } k, \\ 0 & \text{otherwise.} \end{cases}$$

Given the node cluster memberships, the probability of connection is given by:

$$A_{ij} \mid \{\eta_{ik} = 1, \eta_{jq} = 1\} \sim \mathcal{B}(\boldsymbol{\Pi}_{kq}).$$

²[holland1983stochastic](#); [nowicki2001estimation](#); [daudin2006](#).

³[hoff2007modeling](#); [daudin2010model](#).

Stochastic Block Model

The stochastic block model² assumes that each node is assigned to a single cluster:

$$\boldsymbol{\eta}_i \stackrel{i.i.d.}{\sim} \text{Multinomial}(1; \alpha = (\alpha_1, \dots, \alpha_d)). \quad (1)$$

Hence,

$$\eta_{ik} = \begin{cases} 1 & \text{if } i \text{ is in cluster } k, \\ 0 & \text{otherwise.} \end{cases}$$

Given the node cluster memberships, the probability of connection is given by:

$$A_{ij} \mid \{\eta_{ik} = 1, \eta_{jq} = 1\} \sim \mathcal{B}(\boldsymbol{\Pi}_{kq}) = \mathcal{B}(\boldsymbol{\eta}_i^\top \boldsymbol{\Pi} \boldsymbol{\eta}_j).$$

Can we relax the binary constraint from $\boldsymbol{\eta}_i \in \{0, 1\}^Q$ to $\boldsymbol{\eta}_i \in \Delta_Q$ instead ?³

²holland1983stochastic; nowicki2001estimation; daudin2006.

³hoff2007modeling; daudin2010model.

In this work, we assume that the node cluster membership assignment are not binary but continuous giving rise to the following assumptions:

$$\left. \begin{array}{l} \mathbf{z}_i \stackrel{i.i.d}{\sim} \mathcal{N}_d(0, \mathbf{I}_d) \\ \boldsymbol{\eta}_i = \text{softmax}(\mathbf{z}_i) \end{array} \right\} \text{LogisticNormal distribution}$$

In this work, we assume that the node cluster membership assignment are not binary but continuous giving rise to the following assumptions:

$$\left. \begin{aligned} \mathbf{z}_i &\stackrel{i.i.d}{\sim} \mathcal{N}_d(0, \mathbf{I}_d) \\ \boldsymbol{\eta}_i &= \text{softmax}(\mathbf{z}_i) \end{aligned} \right\} \text{LogisticNormal distribution}$$
$$A_{ij} \mid \{\boldsymbol{\eta}_i, \boldsymbol{\eta}_j\} \sim \mathcal{B}(\boldsymbol{\eta}_i^\top \mathbf{\Pi} \boldsymbol{\eta}_j).$$

In this work, we aimed at maximising the **marginal log-likelihood** given by:

$$\log p(\mathbf{A} \mid \mathbf{\Pi}) = \log \int_{\mathbf{Z}} p(\mathbf{A}, \mathbf{Z} \mid \mathbf{\Pi}) d\mathbf{Z}. \quad (2)$$

The marginal likelihood being intractable, we rely on a variational inference to maximise it. In particular, for any distribution $R(\cdot)$ over the latent variable \mathbf{Z} , the following decomposition holds true:

$$\log p(\mathbf{A} \mid \mathbf{\Pi}) = \mathcal{L}(\mathbf{\Pi}; R) + \text{KL}(R(\cdot) \parallel p(\mathbf{Z} \mid \mathbf{A})),$$

where

$$\mathcal{L}(\mathbf{\Pi}; R) = \mathbb{E}_{R(\mathbf{Z})} \left[\log \frac{p(\mathbf{A}, \mathbf{Z} \mid \mathbf{\Pi})}{R(\mathbf{Z})} \right]. \quad (3)$$

The quantity $\mathcal{L}(\mathbf{\Pi}; R)$ is called the **expected lower bound (ELBO)**.

Inference: variational assumption

Assuming that the variational distribution respect the mean-field hypothesis:

$$R_\phi(\mathbf{Z}) = \prod_{i=1}^N \mathcal{N}_d(\mu_\phi(\mathbf{A})_i, \sigma_\phi(\mathbf{A})_i^2 \mathbf{I}_d), \quad (4)$$

where the parameters are the output of a graph convolutional network⁴:

$$(\mu_\phi(\mathbf{A}), \log \sigma_\phi(\mathbf{A})^2) = \text{GCN}_\phi(\mathbf{A}). \quad (5)$$

⁴kipf2016variational.

Hence, the ELBO can be written as

$$\begin{aligned}\mathcal{L}(\mathbf{\Pi}; R_\phi) &= \sum_{j < i} \mathbb{E}_{R_\phi(\mathbf{z})} [\log p(A_{ij} \mid \boldsymbol{\eta}_i, \boldsymbol{\eta}_j, \mathbf{\Pi})] - \sum_{i=1}^N \text{KL}(R(\mathbf{z}_i) \mid p(\mathbf{z}_i)) \\ &= \sum_{j < i} A_{ij} \mathbb{E}_{R_\phi(\mathbf{z})} [\log(\boldsymbol{\eta}_i^\top \mathbf{\Pi} \boldsymbol{\eta}_j)] + (1 - A_{ij}) \mathbb{E}_{R_\phi(\mathbf{z})} [\log(1 - \boldsymbol{\eta}_i^\top \mathbf{\Pi} \boldsymbol{\eta}_j)] \\ &\quad - \sum_{i=1}^N \frac{1}{2} (d\sigma_\phi(\mathbf{A})_i^2 + \|\mu_\phi(\mathbf{A})_i\|_2^2 - d \log \sigma_\phi(\mathbf{A})_i^2 - d) .\end{aligned}\tag{6}$$

Next step: maximisation of $\mathcal{L}(\mathbf{\Pi}; R_\phi)$ with respect to $\mathbf{\Pi}$ and ϕ . We can directly optimise the previous quantity with respect to $\mathbf{\Pi}$ with gradient-based algorithm ... but not with respect to ϕ . Do you see the issue ?

The reparametrisation trick⁵

How to compute the gradient $\frac{\partial}{\partial \phi} \mathcal{L}(\mathbf{\Pi}; R_\phi)$? Based on the previous slide, we have:

$$\frac{\partial}{\partial \phi} \mathcal{L}(\mathbf{\Pi}; R_\phi) = \sum_{j < i} \frac{\partial}{\partial \phi} \mathbb{E}_{R_\phi(\mathbf{z})} [\log p(\textcolor{brown}{A}_{ij} \mid \boldsymbol{\eta}_i, \boldsymbol{\eta}_j, \mathbf{\Pi})] - \sum_{i=1}^N \frac{\partial}{\partial \phi} \overbrace{\text{KL}(R_\phi(\mathbf{z}_i) \mid p(\mathbf{z}_i))}^{\text{analytical form}}. \quad (7)$$

Issue: Since $R_\phi(\cdot)$ depends on ϕ , we cannot interchange the derivative and the integral in the term on the left-hand side.

⁵kingma2014autoencoding; rezende2014stochastic.

The reparametrisation trick⁵

How to compute the gradient $\frac{\partial}{\partial \phi} \mathcal{L}(\mathbf{\Pi}; R_\phi)$? Based on the previous slide, we have:

$$\frac{\partial}{\partial \phi} \mathcal{L}(\mathbf{\Pi}; R_\phi) = \sum_{j < i} \frac{\partial}{\partial \phi} \mathbb{E}_{R_\phi(\mathbf{z})} [\log p(\textcolor{brown}{A}_{ij} \mid \boldsymbol{\eta}_i, \boldsymbol{\eta}_j, \mathbf{\Pi})] - \sum_{i=1}^N \frac{\partial}{\partial \phi} \overbrace{\text{KL}(R_\phi(\mathbf{z}_i) \mid p(\mathbf{z}_i))}^{\text{analytical form}}. \quad (7)$$

Issue: Since $R_\phi(\cdot)$ depends on ϕ , we cannot interchange the derivative and the integral in the term on the left-hand side.

The **reparametrisation trick**⁵ removes this dependency by the following sampling scheme:

$$\epsilon \sim \mathcal{N}_d(0, \mathbf{I}_d), \quad \text{and} \quad \mathbf{z}_i = \mu_\phi(\textcolor{brown}{A})_i + \sigma_\phi(\textcolor{brown}{A})_i \epsilon.$$

Hence, we can now interchange the integral and the derivative and use a Monte-Carlo estimate of the term on the right-hand side of the following equation:

$$\frac{\partial}{\partial \phi} \mathbb{E}_{R_\phi(\mathbf{z})} [\log p(\textcolor{brown}{A}_{ij} \mid \boldsymbol{\eta}_i, \boldsymbol{\eta}_j, \mathbf{\Pi})] = \frac{\partial}{\partial \phi} \mathbb{E}_\epsilon [\log p(\textcolor{brown}{A}_{ij} \mid \boldsymbol{\eta}_i, \boldsymbol{\eta}_j, \mathbf{\Pi})] = \mathbb{E}_\epsilon \left[\frac{\partial}{\partial \phi} \log p(\textcolor{brown}{A}_{ij} \mid \boldsymbol{\eta}_i, \boldsymbol{\eta}_j, \mathbf{\Pi}) \right].$$

⁵kingma2014autoencoding; rezende2014stochastic.

Simulation setup

- ▶ Number of clusters = 5
- ▶ Number of nodes is set to 200
- ▶ $\beta \in \{0.1, 0.2, 0.3\}$ tunes for the high level of connectivity between clusters
- ▶ $\eta = 0.01$ in all our experiments

$$\mathbf{\Pi}^* = \begin{matrix} & \text{Communities} & & \text{Disassortative} & & \text{Hub} \\ \begin{pmatrix} \beta & \epsilon & \dots & \dots & \epsilon \\ \epsilon & \beta & \epsilon & \dots & \epsilon \\ \vdots & \epsilon & \beta & \dots & \epsilon \\ \epsilon & \dots & \beta & \epsilon & \epsilon \\ \epsilon & \epsilon & \dots & \dots & \beta \end{pmatrix} & \begin{pmatrix} \epsilon & \beta & \dots & \dots & \beta \\ \beta & \epsilon & \beta & \dots & \beta \\ \beta & \beta & \epsilon & \dots & \beta \\ \beta & \beta & \dots & \epsilon & \beta \\ \beta & \beta & \dots & \dots & \epsilon \end{pmatrix} & \begin{pmatrix} \beta & \beta & \dots & \dots & \beta \\ \beta & \beta & \epsilon & \dots & \epsilon \\ \beta & \epsilon & \beta & \dots & \epsilon \\ \beta & \epsilon & \dots & \beta & \epsilon \\ \beta & \epsilon & \dots & \dots & \beta \end{pmatrix} \end{matrix}$$

Table: ARI between the true node groups and the estimated ones. The closer to 1 the better.

		Communities	Disassortative	Hub
$\beta = 0.1$	SBM	0.13 ± 0.12	0.03 ± 0.03	0.41 ± 0.14
	ARGVA	0.42 ± 0.08	0.00 ± 0.01	0.07 ± 0.03
	VGAE	0.55 ± 0.11	0.01 ± 0.00	0.20 ± 0.08
	Deep LPBM	0.46 ± 0.20	0.02 ± 0.02	0.18 ± 0.04
$\beta = 0.2$	SBM	0.71 ± 0.15	0.47 ± 0.29	0.83 ± 0.15
	ARGVA	0.85 ± 0.03	0.01 ± 0.01	0.28 ± 0.06
	VGAE	0.99 ± 0.01	0.00 ± 0.01	0.78 ± 0.15
	Deep LPBM	0.99 ± 0.01	0.39 ± 0.14	0.90 ± 0.05
$\beta = 0.3$	SBM	0.77 ± 0.12	0.81 ± 0.21	0.98 ± 0.07
	ARGVA	0.88 ± 0.03	0.06 ± 0.04	0.56 ± 0.22
	VGAE	1.00 ± 0.00	0.00 ± 0.00	0.93 ± 0.07
	Deep LPBM	1.00 ± 0.00	1.00 ± 0.00	1.00 ± 0.00

Real dataset: the French political blogosphere⁶

- ▶ This dataset is composed of 194 nodes
- ▶ Each node corresponds to a political blog
- ▶ An edge exists between two blogs if a link exists between the two

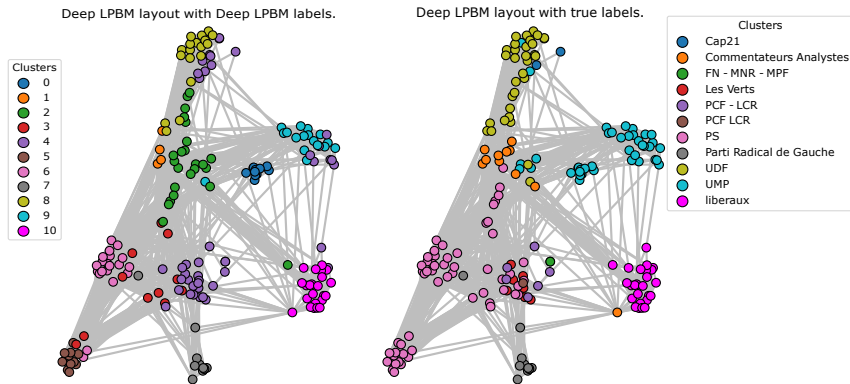


Figure: Representation obtained with Deep LPBM.

Real dataset: the French political blogosphere⁷

- ▶ the blogs related to the Socialist party are split in two clusters (5 and 6) revealing two different connectivity patterns




Figure: Connectivity between the estimated clusters and between the blog associated to the political parties.

⁷zanghi2008Fastonline.

Conclusion

- ▶ The combination of graph neural networks with block modelling provides insightful results
- ▶ The model selection working without GNN still works in the variational autoencoder setting
- ▶ Need to test it on other datasets (in the presence of connectivity patterns different from communities)

Conclusion

- ▶ The combination of graph neural networks with block modelling provides insightful results
- ▶ The model selection working without GNN still works in the variational autoencoder setting
- ▶ Need to test it on other datasets (in the presence of connectivity patterns different from communities)

Thank you for your attention !

Details about VGAE⁸

Denoting $\tilde{\mathbf{A}} = \mathbf{f}^{-1/2}(\mathbf{A} + \mathbf{I}_N)\mathbf{f}^{-1/2}$, the graph convolutional network can be summarised as

⁸kipf2016variational.

Denoting $\tilde{\mathbf{A}} = \mathbf{f}^{-1/2}(\mathbf{A} + \mathbf{I}_N)\mathbf{f}^{-1/2}$, the graph convolutional network can be summarised as

$$\begin{aligned}\mu_\phi(\mathbf{A}) &= \tilde{\mathbf{A}} \operatorname{ReLU}(\tilde{\mathbf{A}}\mathbf{\Omega}_0)\mathbf{\Omega}_\mu, \\ \log \sigma_\phi^2(\mathbf{A}) &= \tilde{\mathbf{A}} \operatorname{ReLU}(\tilde{\mathbf{A}}\mathbf{\Omega}_0)\mathbf{\Omega}_\sigma,\end{aligned}$$

where

- ▶ $\operatorname{ReLU}(x) = (\max(0, x_1), \dots, \max(0, x_F))$ if $x \in \mathbb{R}^F$,
- ▶ $\mathbf{\Omega}_0 \in \mathcal{M}_{N \times D}(\mathbb{R})$ with $D = 64$ in all the experiments we carried out,
- ▶ $\mathbf{\Omega}_\mu, \mathbf{\Omega}_\sigma \in \mathcal{M}_{D \times (Q-1)}(\mathbb{R})$.

⁸kipf2016variational.

In this methodology, we consider the following three model selection criterion:

$$\text{AIC}(Q, \mathcal{M}) = \log p(\mathbf{A} \mid \mathbf{Z}) - \frac{Q * (Q + 1)}{2} - N(Q - 1),$$

$$\text{BIC}(Q, \mathcal{M}) = \log p(\mathbf{A} \mid \mathbf{Z}) - \frac{Q * (Q + 1)/2}{2} \log \left(\frac{N * (N - 1)}{2} \right) - \frac{Q - 1}{2} \log(N),$$

$$\text{ICL}(Q, \mathcal{M}) = \log p(\mathbf{A} \mid \mathbf{Z}) - \frac{Q * (Q + 1)/2}{2} \log \left(\frac{N * (N - 1)}{2} \right) + \log p(\mathbf{Z}) - \frac{Q - 1}{2} \log(N).$$

Model Selection

Table: Comparison of AIC (2a) and BIC (2b) and ICL (2c) to select the best number of clusters for Deep LPBM with $\beta = 0.3$, the true number of cluster corresponds to the shadowed row.

(a) AIC				(b) BIC			(c) ICL		
Q	Commu	Disass	Hub	Commu	Disass	Hub	Commu	Disass	Hub
2	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0
4	0	0	1	0	0	0	0	0	0
5*	10	10	9	0	7	6	0	8	7
6	0	0	0	10	3	4	10	2	3
8	0	0	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0
10	0	0	0	0	0	0	0	0	0
16	0	0	0	0	0	0	0	0	0